

Content-Based Recommendation (Part A)

Topic 5A

Content-Based Recommendation

- ❑ Collaborative recommendation approaches are based on user ratings, without utilizing specific item descriptions.
- ❑ In contrast, content-based recommendations utilize item characteristics (“content”) and user profiles to find recommendable items.
- ❑ A user profile describes the (past) interests of a user in terms of preferred item characteristics.
- ❑ Item descriptions may come from external sources as well as internal processing of item content (characteristics and attributes).

Content Representation and Content Similarity

- When an explicit list of features for each item and user preferences are maintained, the recommendation task is simply consists of **matching item characteristics** and **a user profile of preferences**.

| Title | Genre | Author | Type | Price | Keywords |
|-----------------------------|------------------|----------------|-----------|-------|--|
| <i>The Night of the Gun</i> | Memoir | David Carr | Paperback | 29.90 | press and journalism, drug addiction, personal memoirs, New York |
| <i>The Lace Reader</i> | Fiction, Mystery | Brunonia Barry | Hardcover | 49.90 | American contemporary fiction, detective, historical |

Item Profile and User's Preference Profile

| Title | Genre | Author | Type | Price | Keywords |
|-----------------------------|------------------|----------------|-----------|-------|--|
| <i>The Night of the Gun</i> | Memoir | David Carr | Paperback | 29.90 | press and journalism, drug addiction, personal memoirs, New York |
| <i>The Lace Reader</i> | Fiction, Mystery | Brunonia Barry | Hardcover | 49.90 | American contemporary fiction, detective, historical |

| Title | Genre | Author | Type | Price | Keywords |
|-------|-------------------|-----------------------------|-----------|-------|-----------------------------|
| ... | Fiction, Suspense | Brunonia Barry, Ken Follett | Paperback | 25.65 | detective, murder, New York |

Content Representation and Content Similarity

- User profile can be constructed in various ways
 - ▣ Direct response to questions
 - ▣ Rating of a set of items along different dimensions
 - ▣ Automatic derivation of a set of keywords
 - Dice coefficient measures the similarity between books b_i and b_j as:

$$\frac{2 \times |\text{keywords}(b_i) \cap \text{keywords}(b_j)|}{|\text{keywords}(b_i)| + |\text{keywords}(b_j)|}$$

Content-Based Recommendation Approaches

- The standard approach in content-based recommendation is to use a list of relevant keywords that appear within the document, generated automatically from the document content itself or from a free-text description thereof.
- Boolean vector approach
 - ▣ Problems and limitations
- Term frequency/inverse document frequency (TF-IDF) approach

TF-IDF

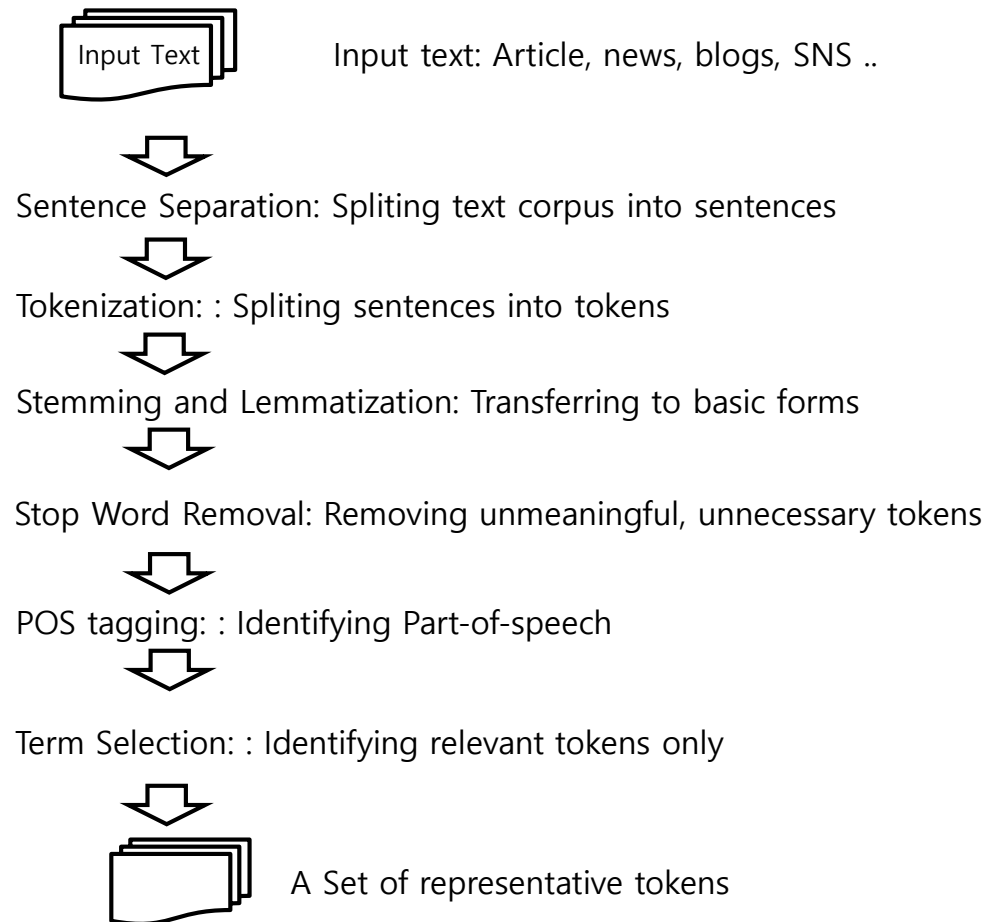
- Assume that N is the total number of documents that can be recommended to users and that keyword k_i appears in n_i of them. Also, assume that $f_{i,j}$ is the number of times keyword k_i appears in the document d_j .
 - $TF_{i,j} = f_{i,j} / \max_z f_{z,j}$
 - Where $f_{z,j}$ of all keywords k_z that appear in the document d_j .
 - $IDF_i = \log(N/n_i)$
 - $W_{i,j} = TF_{i,j} \times IDF_i$

Improving the Vector Space Model

- Stop words and stemming
- Size cutoffs
- Phrases
- Limitations
 - ▣ Context
 - ▣ Feature extraction
 - ▣ Quality
 - ▣ Overspecialization


NLP – Overall Process

NLP process



NLP – Overall Process

NLP process



Input text: Article, news, blogs, SNS ..



Sentence Separation: Splitting text corpus into sentences



Tokenization: : Splitting sentences into tokens



Stemming and Lemmatization: Transferring to basic forms



Stop Word Removal: Removing unmeaningful, unnecessary tokens



POS tagging: : Identifying Part-of-speech



Term Selection: : Identifying relevant tokens only



A Set of representative tokens

Working Example

Figure 1 shows the overall architecture of our framework...

Figure/ 1 / shows / the / overall / architecture / of /our/ framework.../

Fig/ 1 / show / the / overall / architecture / of /our/ framework.../

Fig/ show / overall / architecture / framework.../

Fig (NN)/ show(VB) / overall (ADV) / architecture (NN) / framework (NN).../

Fig (NN)/ architecture (NN) / framework (NN).../

Stemming and Lemmatization

- Lemmatization: Reduce inflectional/variant forms to base form
 - E.g.,
am, are, is -> be
car, cars, car's, cars' -> car
 - *the boy's cars are different colors -> the boy car be different color*
 - Lemmatization implies doing “proper” reduction to dictionary headword form
- Stemming: Reduce terms to their “roots” before indexing
 - “Stemming” suggest crude affix chopping
 - e.g., **automate(s), automatic, automation** all reduced to **automat**.
- Note: Stemming would hurt search precision over recall:
 - Stemming returns ‘oper’ from ‘*operate operating operates operation operative operatives operational*’
 - since *operate* in its various forms is a common verb, we would expect to lose considerable precision on queries (i.e., operational and research, operating and system, operative and dentistry)

Useful NLP Tools - Python

- NLTK (Natural Language Toolkit): <http://www.nltk.org/>
 - ▣ used for such tasks as tokenization, lemmatization, stemming, parsing, POS tagging, etc. This library has tools for almost all NLP tasks
- Spacy: <https://spacy.io/>
 - ▣ the main competitor of the NLTK. These two libraries can be used for the same tasks.
- Scikit-learn: <https://scikit-learn.org/>
 - ▣ provides a large library for machine learning. The tools for text preprocessing are also included.
- Gensim: <https://radimrehurek.com/gensim/>
 - ▣ the package for topic and vector space modeling, document similarity.
- KoNLPy: <https://konlpy-ko.readthedocs.io/ko/v0.4.3/>
 - ▣ a Python package for Korean natural language processing

| | ⊕ PROS | ⊖ CONS |
|---|---|--|
|  | <ul style="list-style-type: none"> + The most well-known and full NLP library + Many third-party extensions + Plenty of approaches to each NLP task + Fast sentence tokenization + Supports the largest number of languages compared to other libraries | <ul style="list-style-type: none"> - Complicated to learn and use - Quite slow - In sentence tokenization, NLTK only splits text by sentences, without analyzing the semantic structure - Processes strings which is not very typical for object-oriented language Python - Doesn't provide neural network models - No integrated word vectors |
|  | <ul style="list-style-type: none"> + The fastest NLP framework + Easy to learn and use because it has one single highly optimized tool for each task + Processes objects; more object-oriented, comparing to other libs + Uses neural networks for training some models + Provides built-in word vectors + Active support and development | <ul style="list-style-type: none"> - Lacks flexibility, comparing to NLTK - Sentence tokenization is slower than in NLTK - Doesn't support many languages. There are models only for 7 languages and "multi-language" models |
|  | <ul style="list-style-type: none"> + Has functions which help to use the bag-of-words method of creating features for the text classification problems + Provides a wide variety of algorithms to build machine learning models + Has good documentation and intuitive classes' methods | <ul style="list-style-type: none"> - For more sophisticated preprocessing things (for example, pos-tagging), you should use some other NLP library and only after it you can use models from scikit-learn - Doesn't use neural networks for text preprocessing |
|  | <ul style="list-style-type: none"> + Works with large datasets and processes data streams + Provides tf-idf vectorization, word2vec, document2vec, latent semantic analysis, latent Dirichlet allocation + Supports deep learning | <ul style="list-style-type: none"> - Designed primarily for unsupervised text modeling - Doesn't have enough tools to provide full NLP pipeline, so should be used with some other library (Spacy or NLTK) |

Text Preprocessing in Python

□ Helpful Links

▣ Text Preprocessing in Python: Steps, Tools, and Examples

■ <https://medium.com/@datamonsters/text-preprocessing-in-python-steps-tools-and-examples-bf025f872908>

▣ How to Clean Text for Machine Learning with Python

■ <https://machinelearningmastery.com/clean-text-machine-learning-python/>

▣ How to Implement TF-IDF

■ <https://medium.freecodecamp.org/how-to-process-textual-data-using-tf-idf-in-python-cd2bbc0a94a3>

Useful NLP Tools - Java

- **Stanford NLP:** <https://nlp.stanford.edu/software>
 - ▣ Well-known, widely used for English NLP tasks
 - ▣ Offers many features, but unreliable for Korean preprocessing
- **Berkeley NLP:** <http://nlp.cs.berkeley.edu/software.shtml>
 - ▣ Offers a smaller set of functionalities, but its parser works well
- **Apache OpenNLP:** <http://opennlp.apache.org/>
 - ▣ Supports common NLP tasks
 - ▣ Source code is offered as open source
- **Komoran:** <https://www.shineware.co.kr/products/komoran/>
 - ▣ Reliable performance for Korean
 - ▣ Can be easily used with open source search library such as Lucene, Indri
 - ▣ Delay for initial loading

Stanford Core NLP Suite

□ *Input text: “Annie has a little lamb. She is very cute.”*

| Function | Output |
|-----------|--|
| ssplit | “Annie has a little lamb.” “She is very cute.” |
| token | Annie, has, a, little, lamb, She, is, very, cute |
| pos | Annie/NNP, has/VBZ, a/DT, little/JJ, lamb/NN, She/PRP, is/VBZ, very/RB, cute/JJ |
| lemma | Annie, have, a, little, lamb, she, be, very, cute |
| ner | She/PERSON, has/O, a/O, little/O, lamb/O, She/O, is/O, very/O, cute/O |
| sentiment | “Annie has a little lamb”/Negative “She is very cute”/Positive |

Similarity-Based Retrieval

□ Nearest neighbors

▣ To recommend documents, we need

- Some history of like/dislike statements made by the user about previous items
- Similarity measure – cosine similarity

▣ K-nearest-neighbor method (KNN)

- Varies the neighborhood size k
- More weight on keywords associated with recent ratings
- Long-term vs. short-term interests

Relevance Feedback – Rocchio

Method

- User provides feedback on the relevance of documents retrieved so as to improve retrieval results in the next round.
- The Rocchio algorithm splits the already rated documents into two groups, D^+ and D^- , and calculate a prototype (or average) vector for these categories. The current query Q_i is then repeatedly refined to Q_{i+1} as follows:

$$Q_{i+1} = \alpha * Q_i + \beta \left(\frac{1}{|D^+|} \sum_{d^+ \in D^+} d^+ \right) - \gamma \left(\frac{1}{|D^-|} \sum_{d^- \in D^-} d^- \right) \quad (3.5)$$

Rocchio Method Formula

- In the formula, the value of α describes how strongly the last query should be weighted while β and γ are control parameters that are used to set the relative importance of positive and negative examples.
- For instance, if $\alpha = 1.5$, $\beta = 2$, and $\gamma = 1$, we don't want the negative examples to have as strong influence as the positive examples.

Rocchio Method: Example


- Let's assume that we have identified 118 terms from the set of documents classified into the categories of medicine, energy, and environment.
 - ▣ The weight of each term represents the importance of the respective term for the category
 - ▣ What is the weight of term 'nuclear' in the category 'medicine'?
- POS_{medicine} contains the documents Doc1-Doc4, and NEG_{medicine} contains the documents Doc5-Doc10
 - $|D^+| = 4$ and $|D^-| = 6$.

Rocchio Method: Example

- Weights of term 'nuclear' in documents in POS_{medicine}
 - ▣ $w_{\text{nuclear_doc1}} = 0.5$
 - ▣ $w_{\text{nuclear_doc2}} = 0$
 - ▣ $w_{\text{nuclear_doc3}} = 0$
 - ▣ $w_{\text{nuclear_doc4}} = 0.5$
- Weight in documents in NEG_{medicine}
 - ▣ $w_{\text{nuclear_doc6}} = 0.5$
- Weight of 'nuclear' in the category 'medicine':
 - ▣ $2 * (0.5 + 0.5) / 4 - 1 * 0.5 / 6 = 0.5 - 0.08 = 0.42$

Relevance Feedback Effect

Figure 3.2.



After feedback, the original query is moved toward the cluster of the relevant documents.