# Data Augmentation & Transfer Learning

regularization methods specialized for CNN

- http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture7.pdf
- http://cs231n.stanford.edu/slides/2018/cs231n_2018_lecture13.pdf

# Data Augmentation

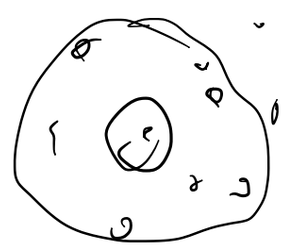

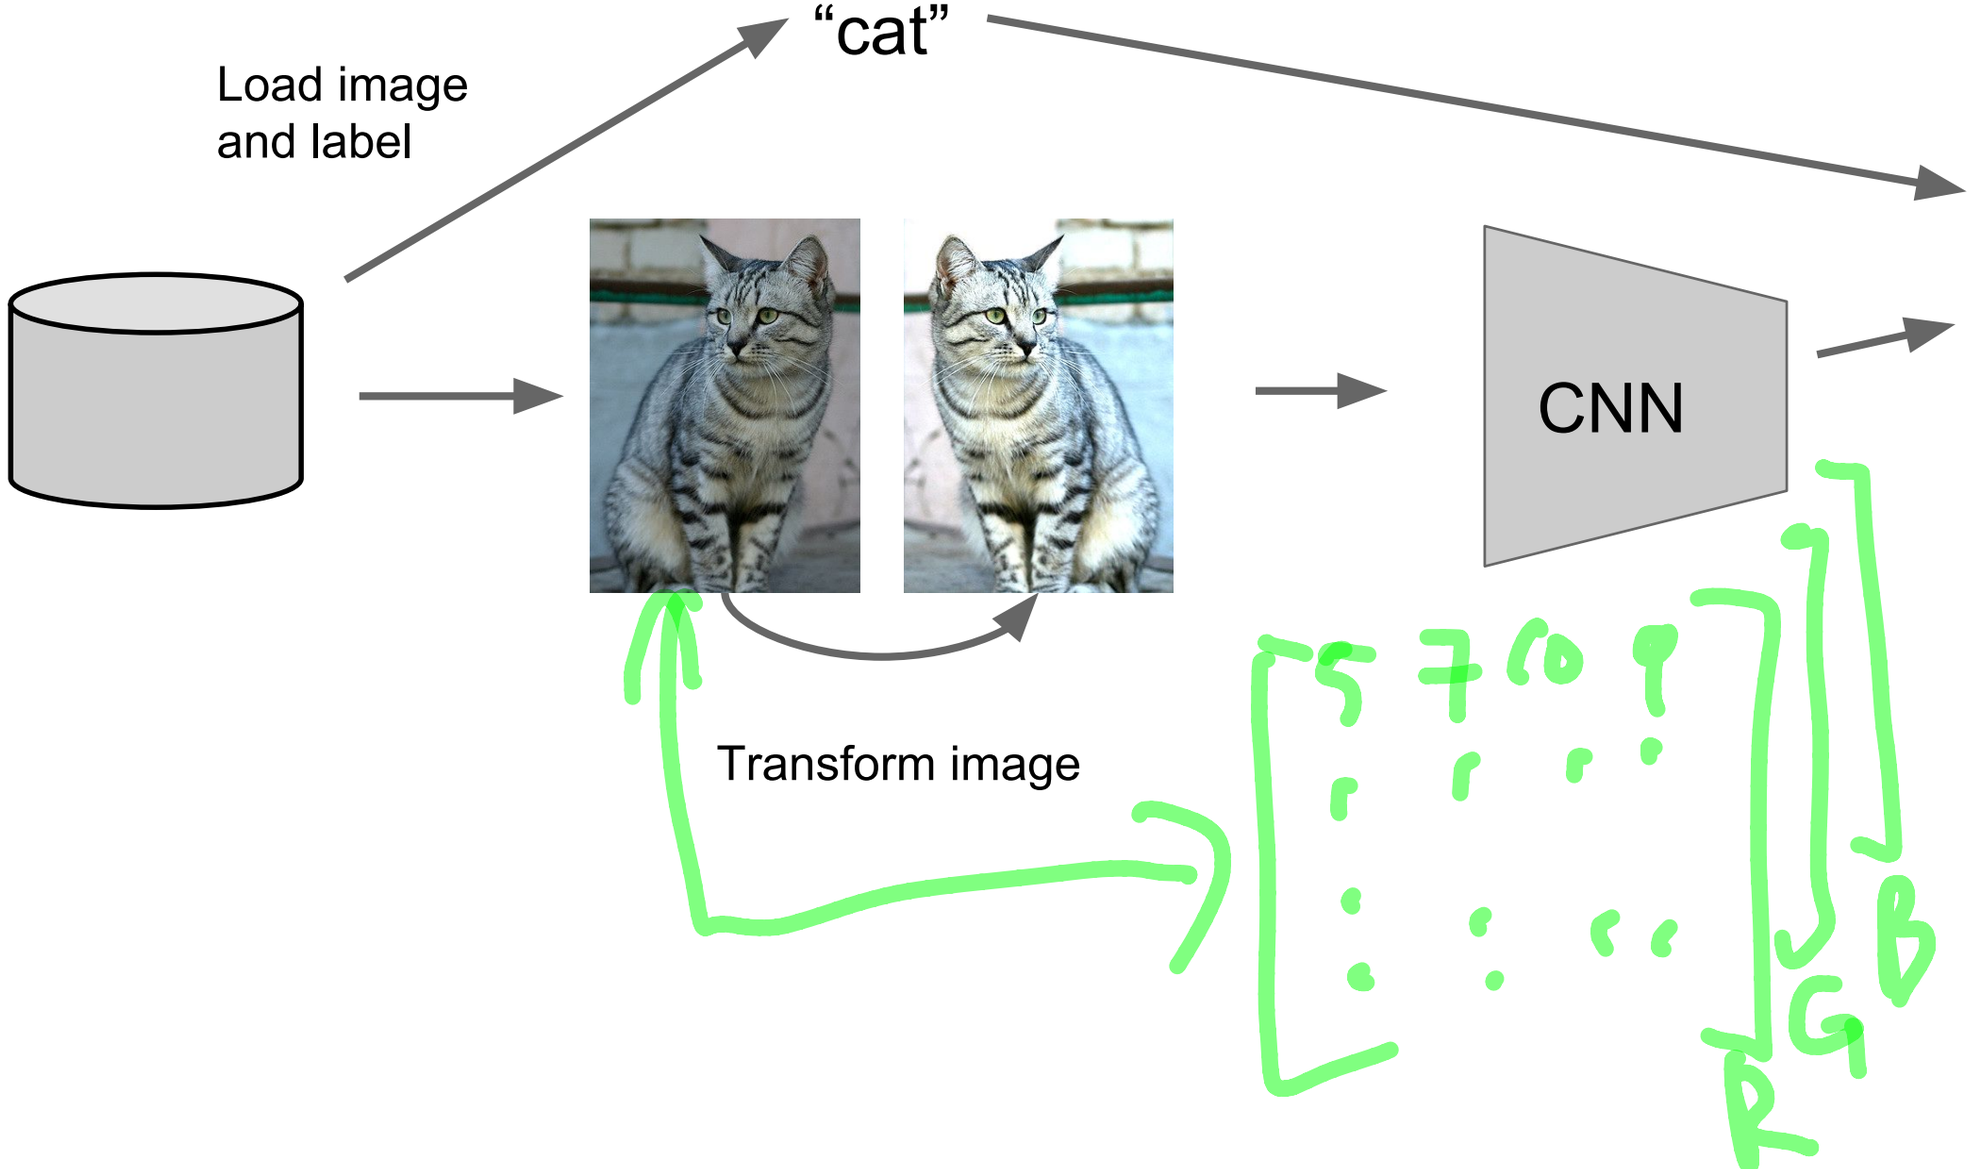

# Data Augmentation

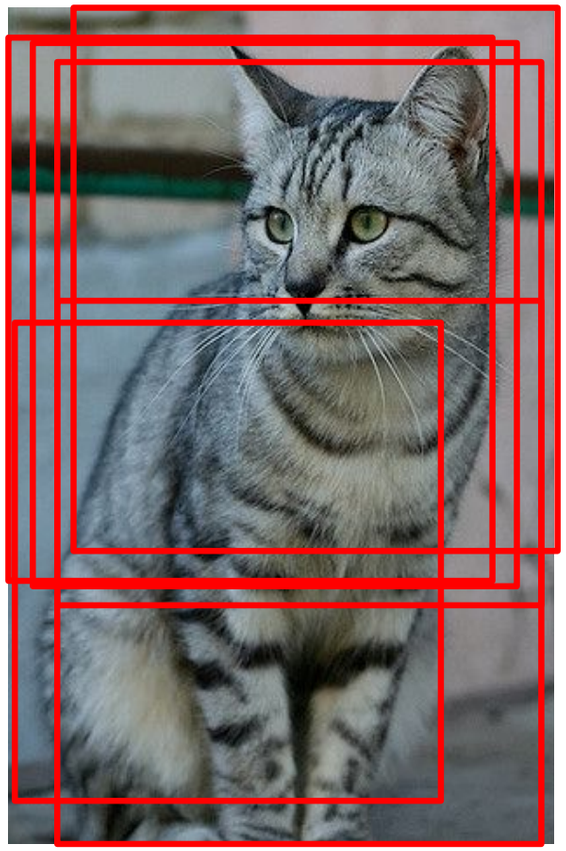

**crop/scale**

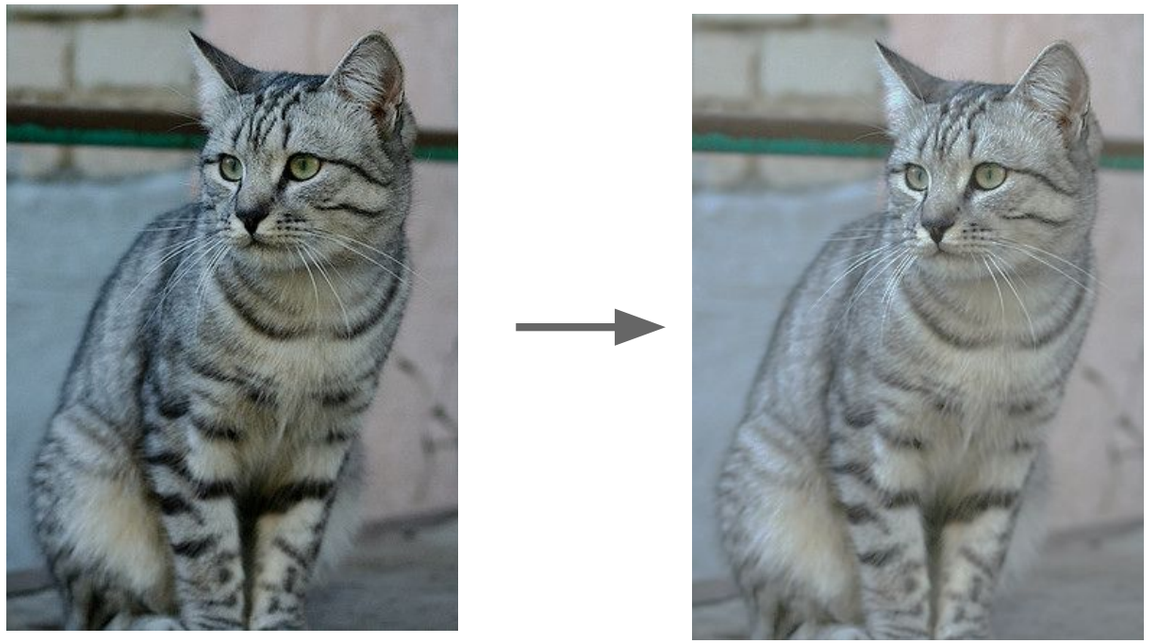

**color jitter**

# Transfer Learning with CNNs

Donahue et al, “DeCAF: A Deep Convolutional Activation Feature for Generic Visual Recognition”, ICML 2014
Razavian et al, “CNN Features Off-the-Shelf: An Astounding Baseline for Recognition”, CVPR Workshops 2014

## 1. Train on Imagenet




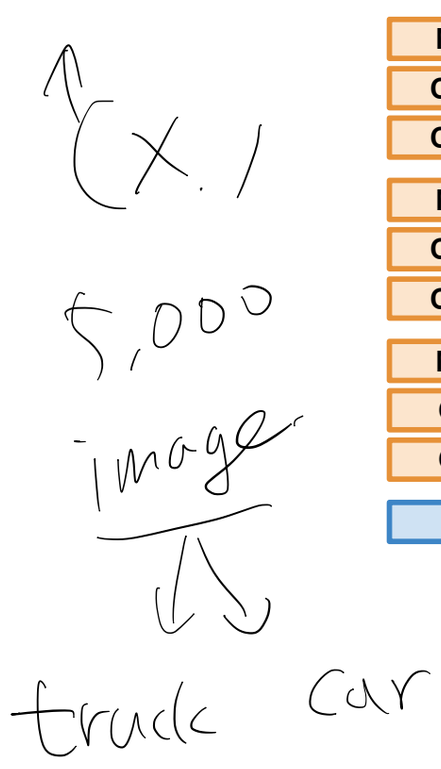


## 2. Small Dataset (C classes)

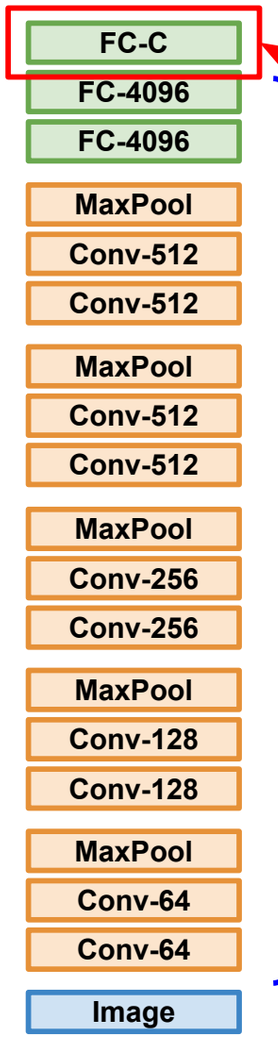


Reinitialize this and train

Freeze these

## 3. Bigger dataset


Train these

With bigger dataset, train more layers

Freeze these

Lower learning rate when finetuning; 1/10 of original LR is good starting point

# Transfer Learning

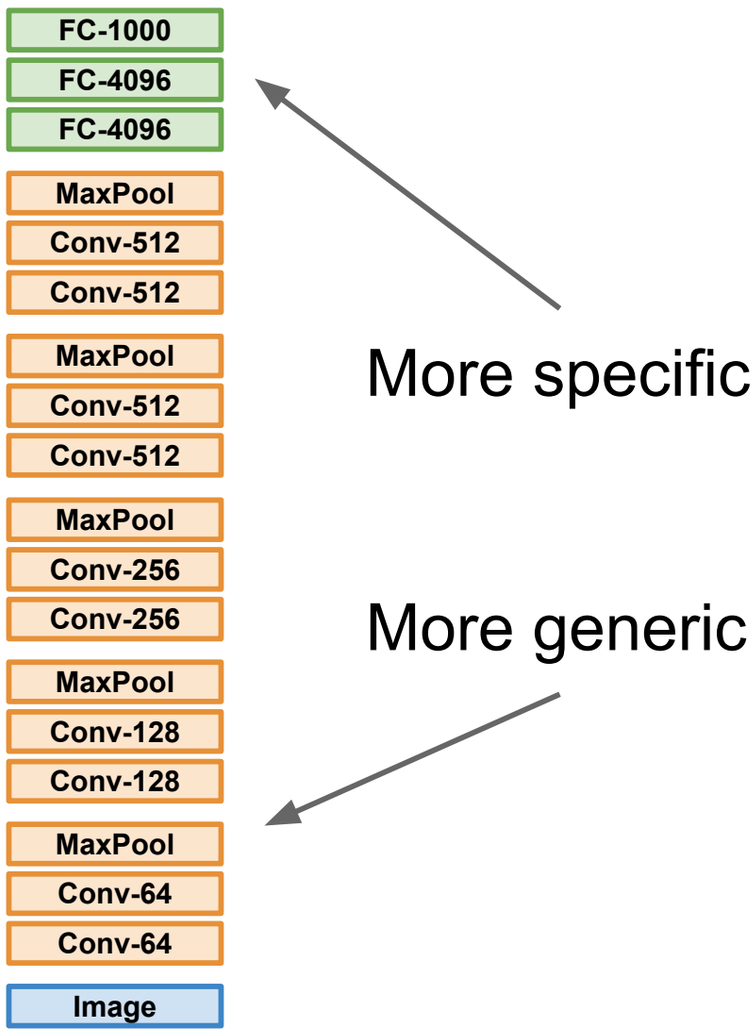


| | **very similar dataset** | **very different dataset** |
|---|---|---|
| **very little data** | Use Linear Classifier on top layer | You're in trouble… Try linear classifier from different stages |
| **quite a lot of data** | Finetune a few layers | Finetune a larger number of layers |

# Transfer Learning

## Transfer learning with CNNs is pervasive…
(it's the norm, not an exception)

Object Detection
(Fast R-CNN)

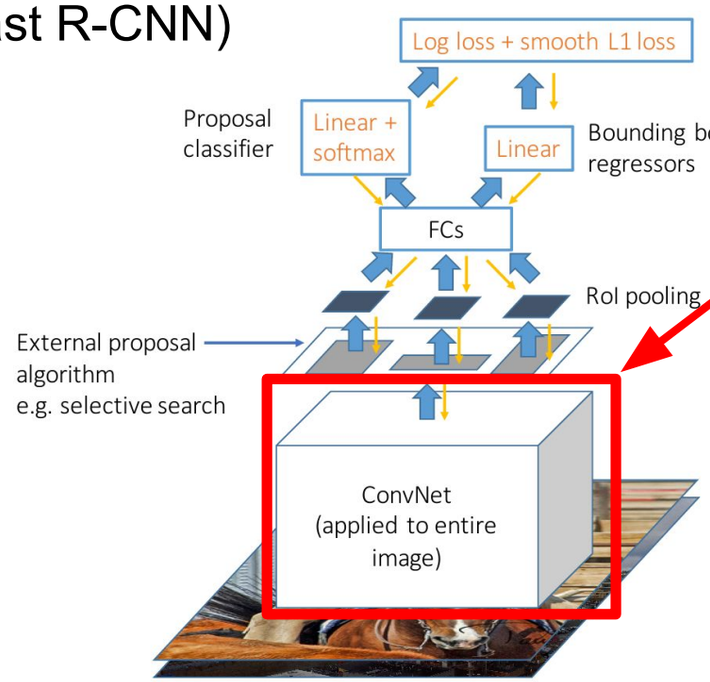


CNN pretrained on ImageNet

Image Captioning: CNN + RNN

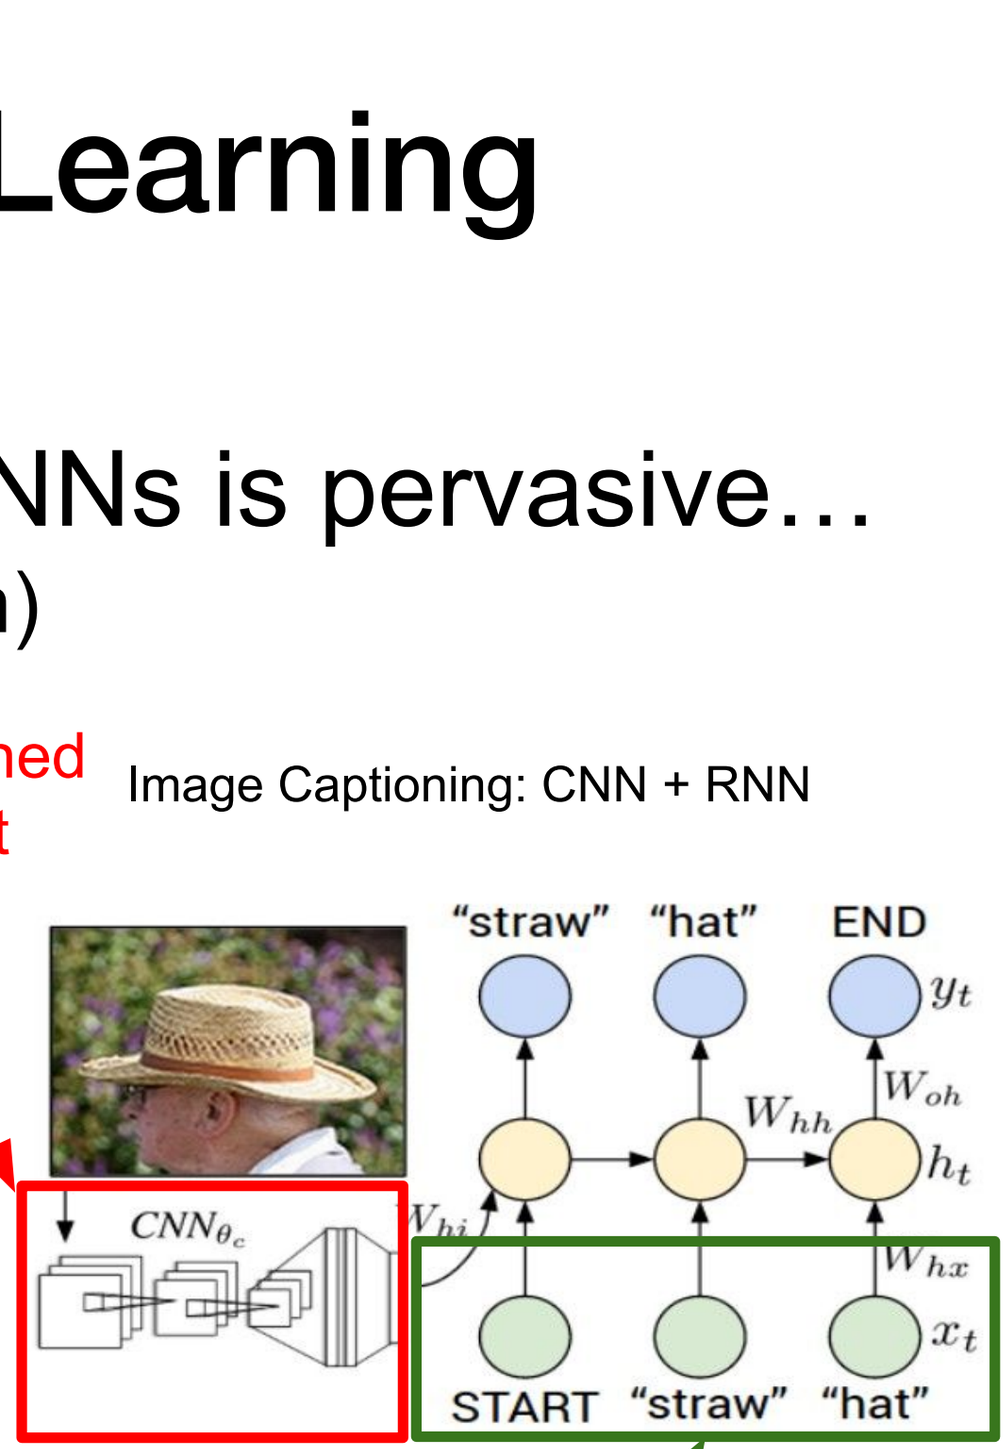


Word vectors pretrained with word2vec

Girshick, "Fast R-CNN", ICCV 2015
Figure copyright Ross Girshick, 2015. Reproduced with permission.

Karpathy and Fei-Fei, "Deep Visual-Semantic Alignments for Generating Image Descriptions", CVPR 2015
Figure copyright IEEE, 2015. Reproduced for educational purposes.

# Style Transfer

Style image

Output image

Content image

Style Target

$y_s$

$\hat{y}$

$y_c$

Content Target

$\ell^{\phi,\texttt{relu1\_2}}_{style}$ $\ell^{\phi,\texttt{relu2\_2}}_{style}$ $\ell^{\phi,\texttt{relu3\_3}}_{style}$ $\ell^{\phi,\texttt{relu4\_3}}_{style}$

Loss Network $\phi$

$\ell^{\phi,\texttt{relu3\_3}}_{feat}$

Gatys, Ecker, and Bethge, “Image style transfer using convolutional neural networks”, CVPR 2016
Figure adapted from Johnson, Alahi, and Fei-Fei, “Perceptual Losses for Real-Time Style Transfer and Super-Resolution”, ECCV 2016. Copyright Springer, 2016. Reproduced for educational purposes.

# What's going on inside ConvNets?



Input Image:
3 x 224 x 224

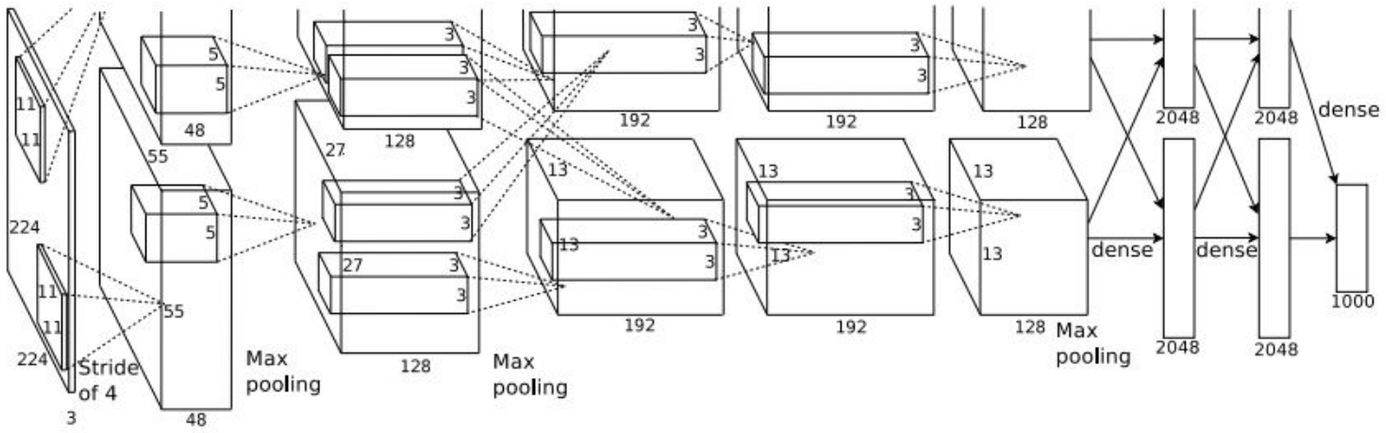


Class Scores:
1000 numbers

What are the intermediate features looking for?

# Example: First layer

## First Layer: Visualize Filters

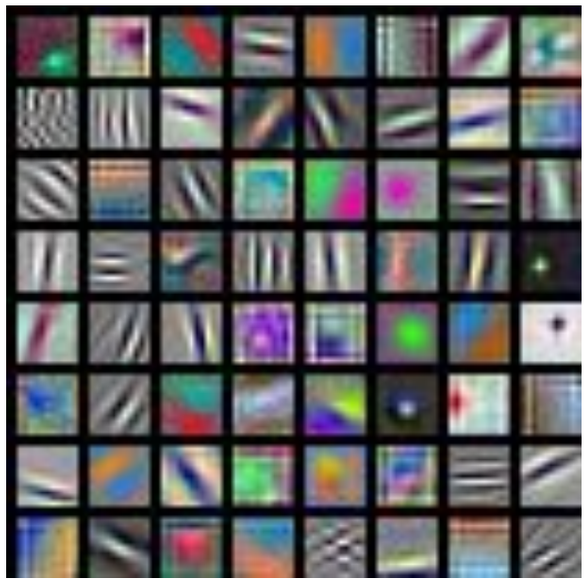

AlexNet:
64 x 3 x 11 x 11

ResNet-18:
64 x 3 x 7 x 7

ResNet-101:
64 x 3 x 7 x 7

DenseNet-121:
64 x 3 x 7 x 7

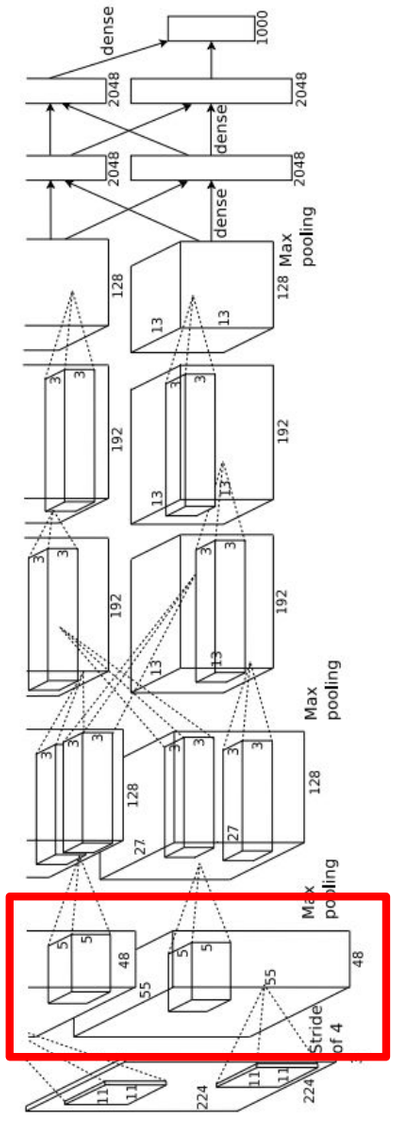


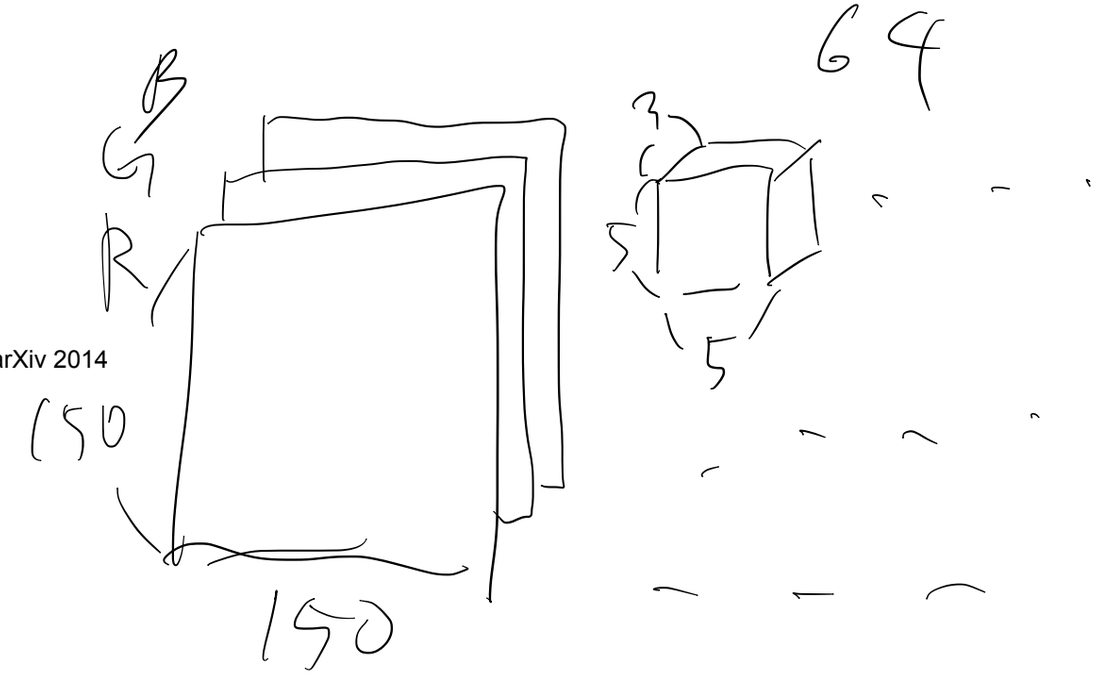


Krizhevsky, “One weird trick for parallelizing convolutional neural networks”, arXiv 2014
He et al, “Deep Residual Learning for Image Recognition”, CVPR 2016
Huang et al, “Densely Connected Convolutional Networks”, CVPR 2017

# Example: Last layer

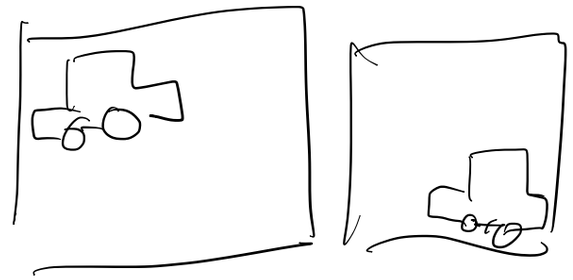

## Last Layer: Nearest Neighbors

4096-dim vector

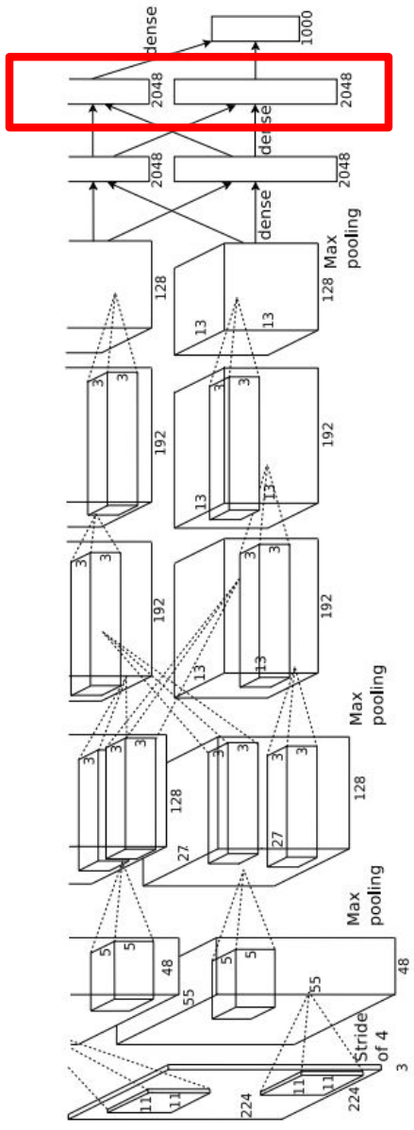


Test image

L2 Nearest neighbors in <u>feature</u> space

**Recall**: Nearest neighbors in <u>pixel</u> space

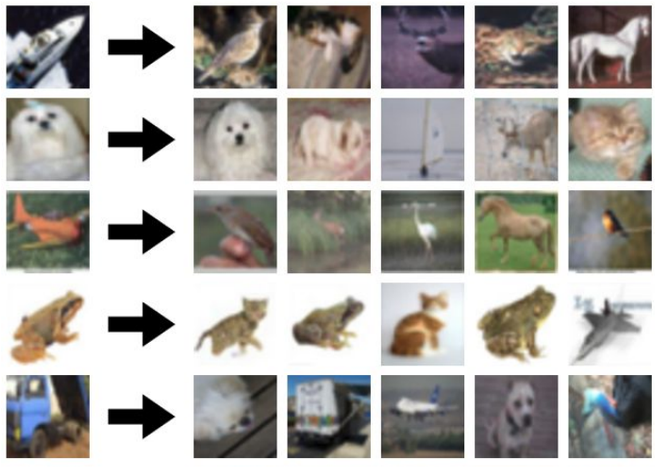

Krizhevsky et al, “ImageNet Classification with Deep Convolutional Neural Networks”, NIPS 2012.
Figures reproduced with permission.

# Example: Last Layer

## Last Layer: Dimensionality Reduction

Visualize the “space” of FC7 feature vectors by reducing dimensionality of vectors from 4096 to 2 dimensions

Simple algorithm: Principal Component Analysis (PCA)

More complex: **t-SNE**

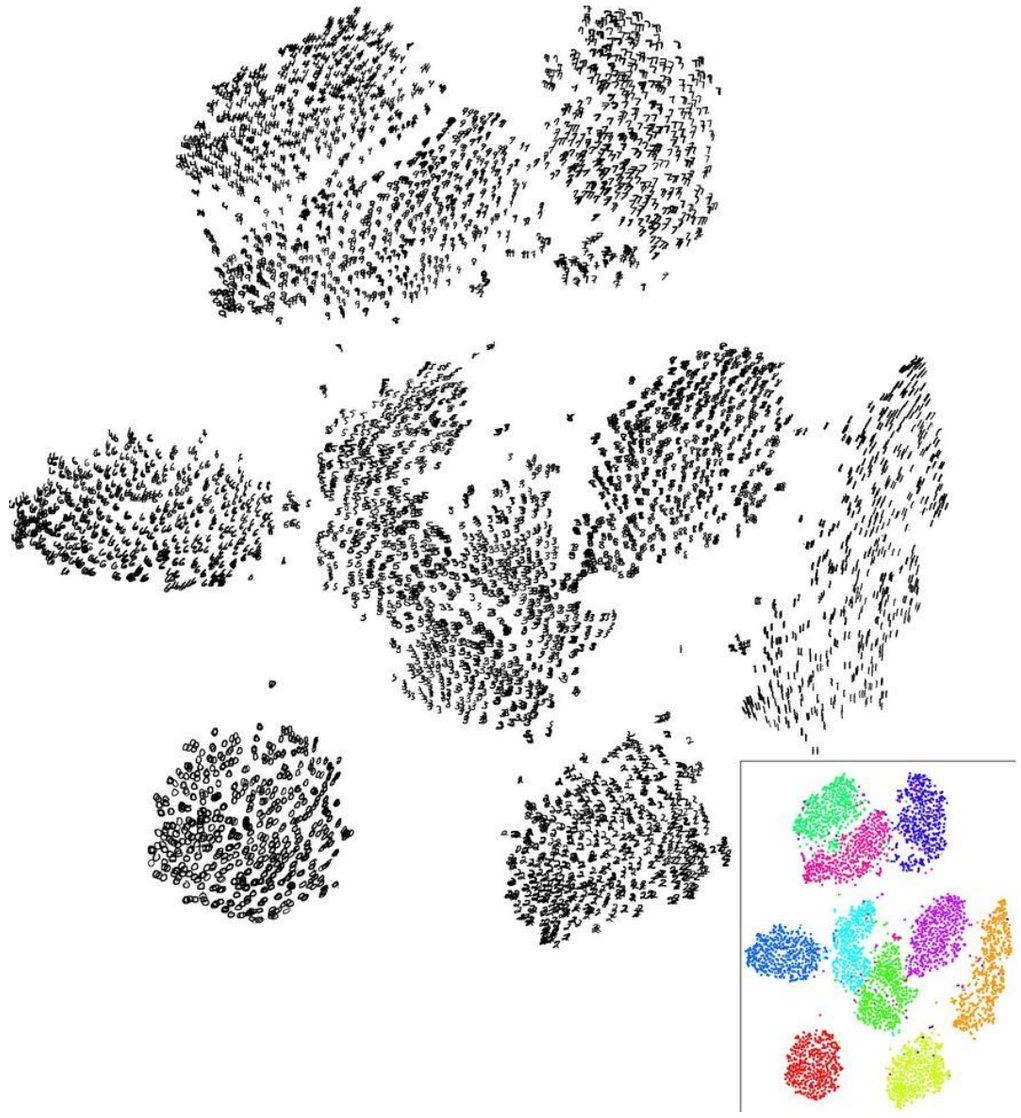

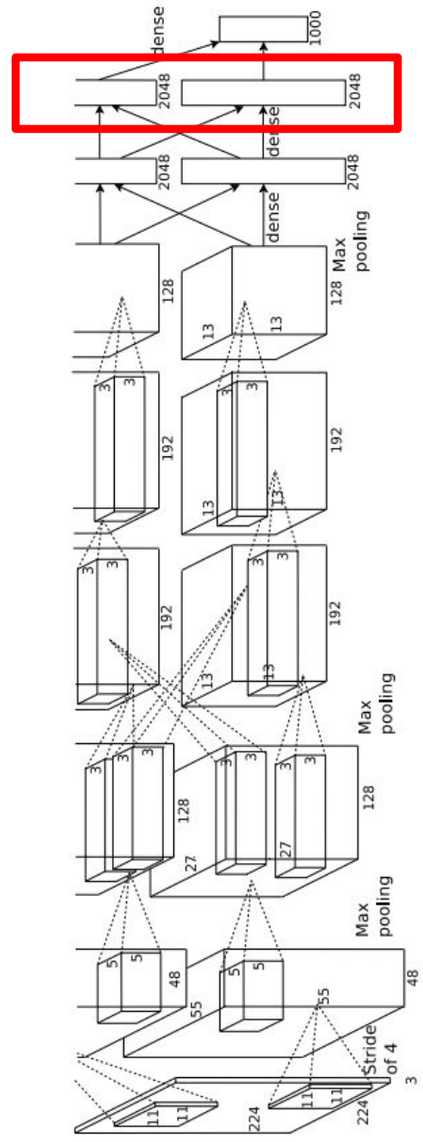


Van der Maaten and Hinton, “Visualizing Data using t-SNE”, JMLR 2008

# Neurons?

## Maximally Activating Patches

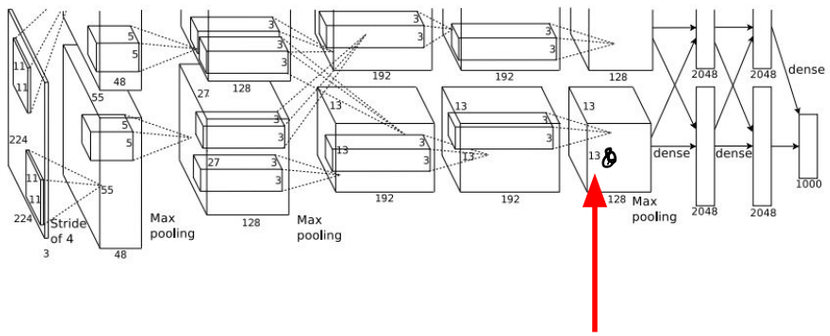


Pick a layer and a channel; e.g. conv5 is 128 x 13 x 13, pick channel 17/128

Run many images through the network, record values of chosen channel

Visualize image patches that correspond to maximal activations

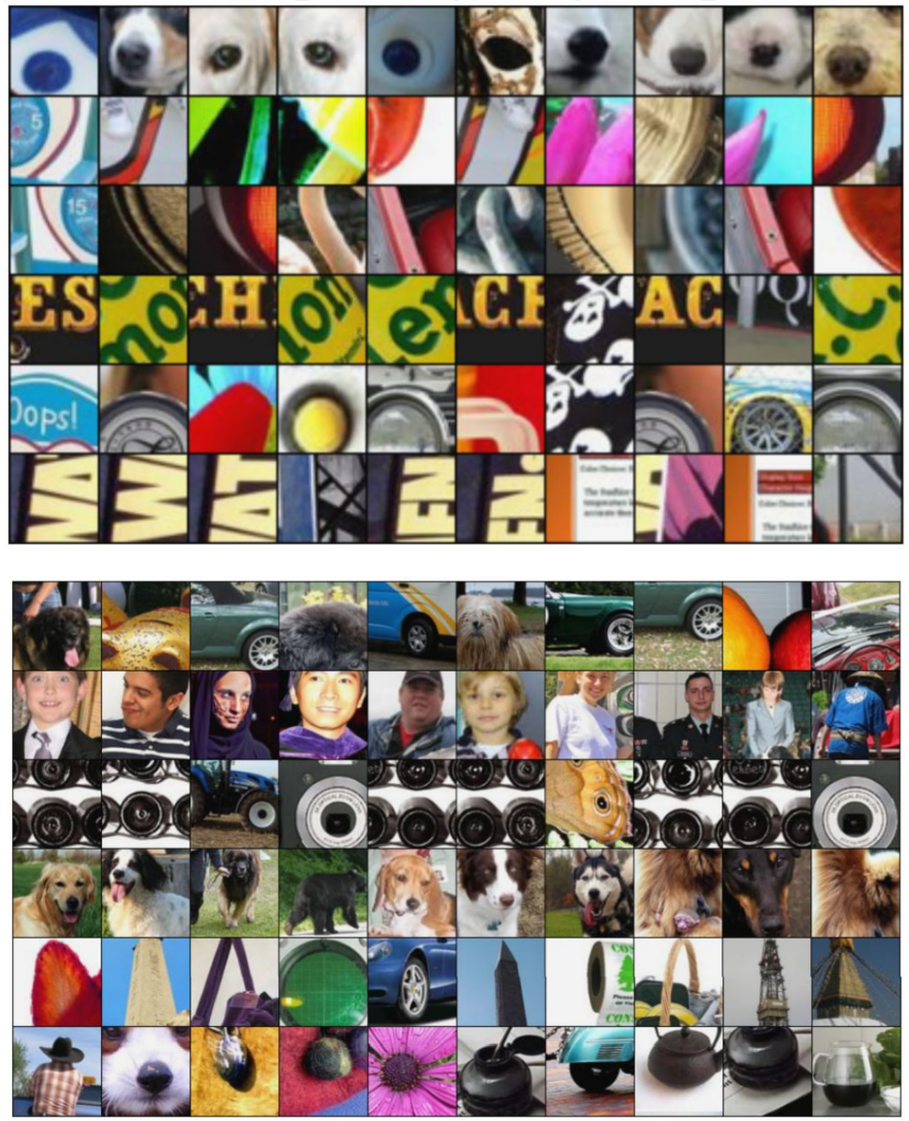

Springenberg et al, “Striving for Simplicity: The All Convolutional Net”, ICLR Workshop 2015
Figure copyright Jost Tobias Springenberg, Alexey Dosovitskiy, Thomas Brox, Martin Riedmiller, 2015; reproduced with permission.