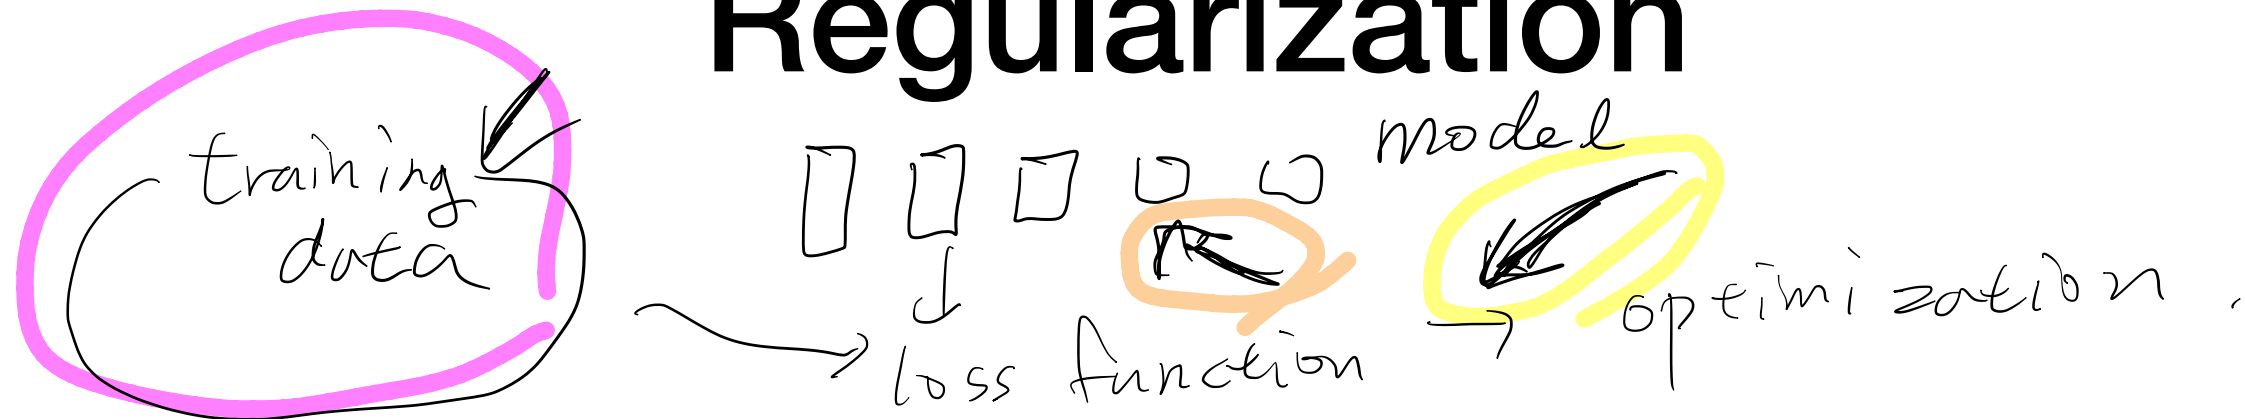


# Regularization

Seyoung Yun

- [http://cs231n.stanford.edu/slides/2017/cs231n\\_2017\\_lecture7.pdf](http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture7.pdf)
- N. Srivastava et al, “Dropout: A Simple Way to Prevent Neural Networks from Overfitting” <http://jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>
- Sergey Ioffe and Christian Szegedy “Batch Normalization: Accelerating Deep Network Training by Reducing Internal Covariate Shift” <https://arxiv.org/abs/1502.03167>
- C. Zhang et al “Understanding deep learning requires rethinking generalization” <https://arxiv.org/abs/1611.03530>


# Regularization



- “A regularizer is anything that hurts the training process”

C. Zhung at ICLR2017 (<https://www.youtube.com/watch?v=kCj51pTQPKI>)

- data augmentation
- weight decay - with an additional cost
- dropout - by adding random noise



# Linear Regression

- RSS: cost of linear regression

$$\mathcal{L}(w, b) = \sum_{i=1}^m (\underbrace{y^{(i)}}_{\text{target}} - \underbrace{w^\top x^{(i)} + b}_{\text{linear function}})^2$$

- regularizer

$$\mathcal{L}(w, b) = \frac{1}{m} \sum_{i=1}^m (y^{(i)} - w^\top x^{(i)} - b)^2 + \frac{\lambda}{2m} \|w\|_2^2$$

or

$$\mathcal{L}(w, b) = \frac{1}{m} \sum_{i=1}^m (y^{(i)} - w^\top x^{(i)} - b)^2 + \frac{\lambda}{2m} \|w\|_1$$

# Weight Decay

$$J(w^{[1]}, b^{[1]}, \dots, w^{[L]}, b^{[L]}) = \frac{1}{m} \sum_{i=1}^m \ell(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{2m} \sum_{i=1}^L \|w^{[i]}\|_F^2$$



$$\|w^{[L]}\|_F^2 = \sum_j \sum_k (w_{jk}^{[L]})^2$$

$$\Rightarrow \nabla_{w^{[L]}} J = \frac{1}{m} \sum_{i=1}^m \nabla_{w^{[L]}} \ell(\hat{y}^{(i)}, y^{(i)}) + \frac{\lambda}{m} \cdot w^{[L]}$$


Back prop.

$$\begin{aligned} w^{[L]}(t+1) &= w^{[L]}(t) - \gamma \cdot \nabla_{w^{[L]}} J(w^{[L]}(t)) \\ &= \left(1 - \frac{\gamma \lambda}{m}\right) w^{[L]}(t) - \underbrace{\frac{\gamma}{m} \sum_{i=1}^m \nabla_{w^{[L]}} \ell(\hat{y}^{(i)}, y^{(i)})}_{\text{Back prop.}} \end{aligned}$$

# Dropout

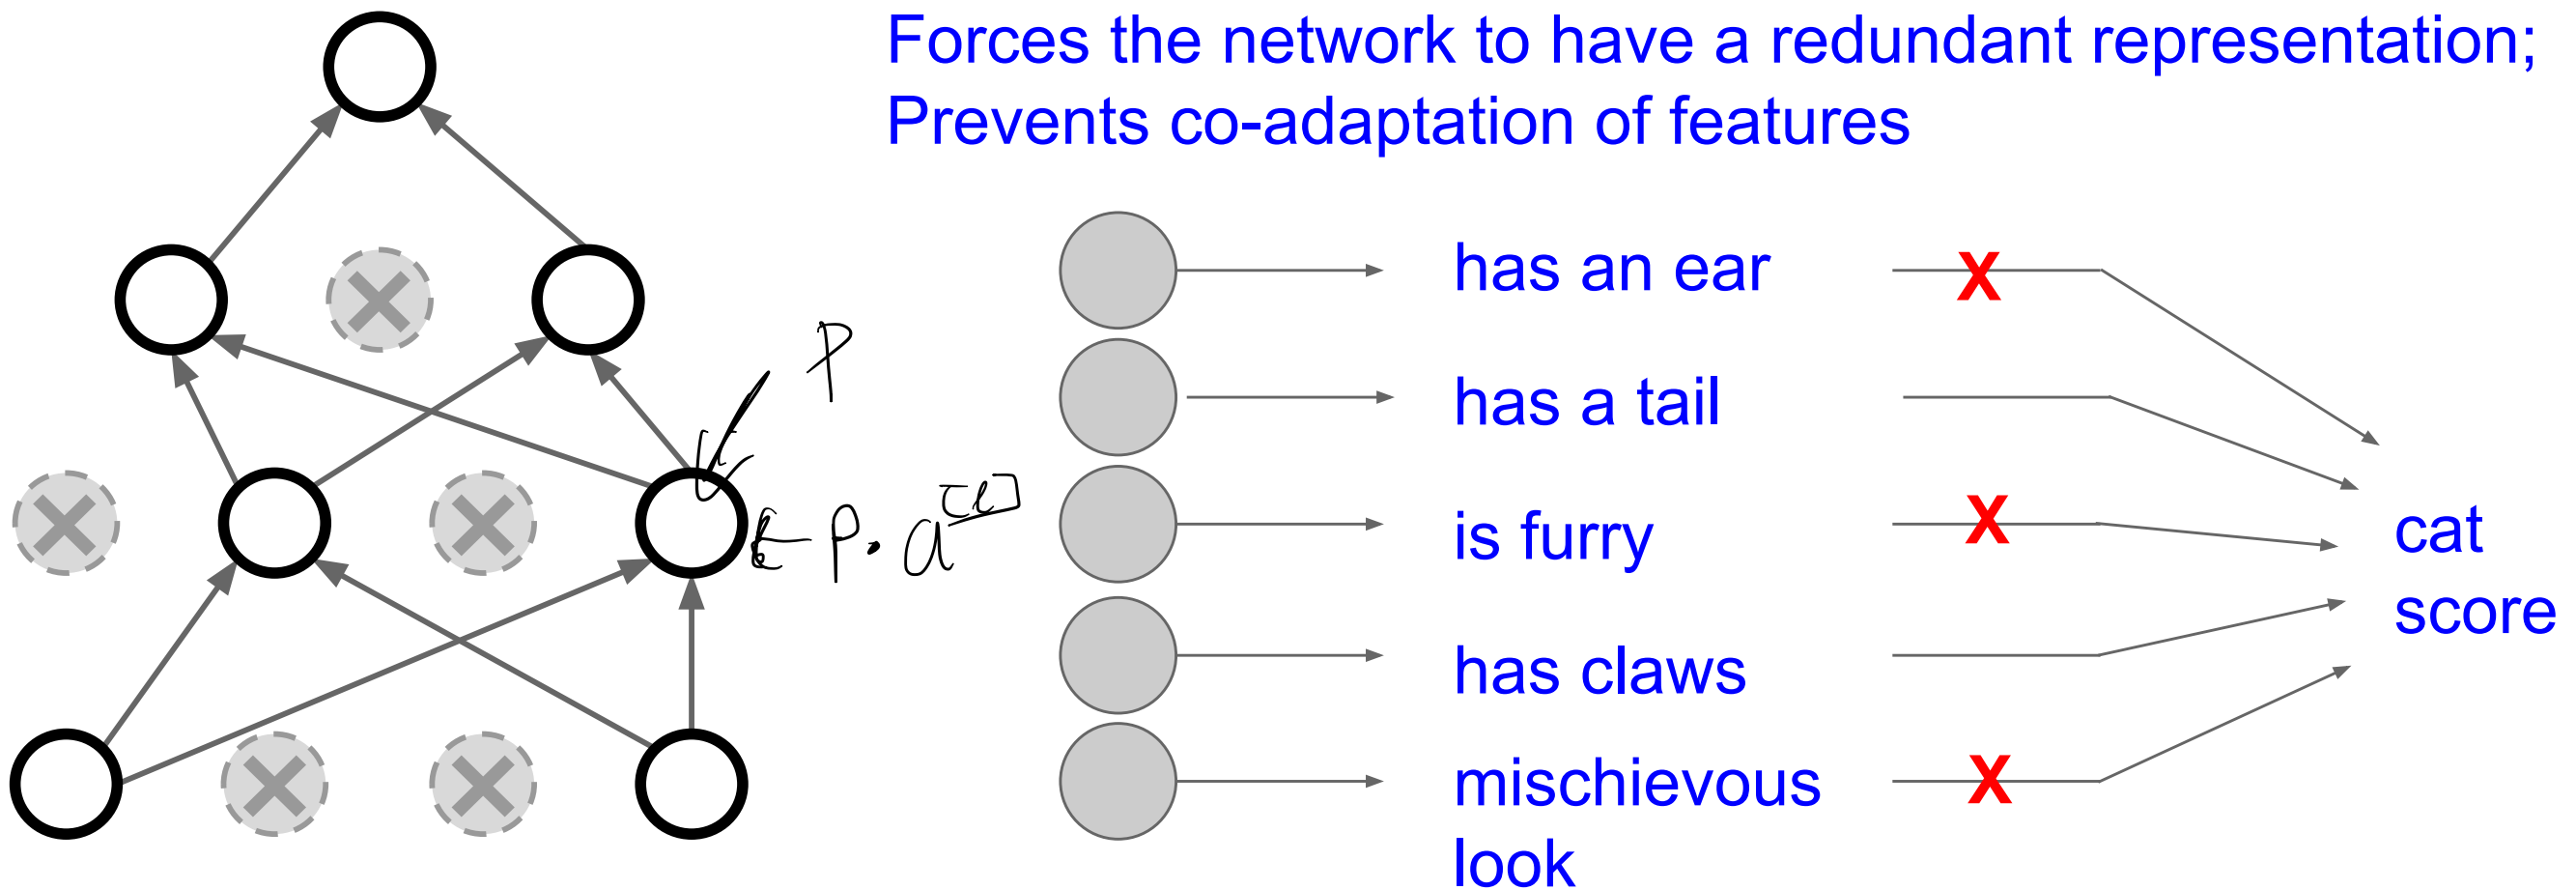


- In each forward pass, randomly erase neurons






# Dropout

- Why can this be good?



# Ensemble of models



- Dropout is training a large ensemble of models (that share parameters).
- Each binary mask is one model



# Dropout: Test time

No dropout -

$$a^{[0]} = X$$

$$z^{[1]} = W^{[1]} a^{[0]} + b^{[1]}$$




$$a^{[1]} = \sigma^{[1]}(z^{[1]})$$

$$\hat{y} = a^{[L]}$$

$$\Rightarrow a^{[L]} = p^{[L]} \cdot \sigma^{[L]}(z^{[L]})$$

- No dropout at test time
- scaling by dropout probability

# Dropout: uncertainty





- Dropout generates ensemble of models
- From the ensemble, estimate mean and variance of the output

# SGD and Early stopping

- SGD adds noises to the network -> a regularizer


- Early stopping

$$\mathcal{L}(w) = \sum_{\tilde{i}=1}^n \ell(y^{(\tilde{i})}, \hat{y}^{(\tilde{i})})$$



# Batch Normalization


- Covariate shift



- “The change in the distributions of layers’ inputs presents a problem because the layers need to continuously adapt to the new distribution. When the input distribution to a learning system changes, it is said to experience covariate shift”


# Batch Normalization

$N(0, 1)$



$$z^{[l+1]} = W^{[l+1]} a^{[l]} + b^{[l+1]}$$

$$a^{[l+1]} = \sigma(z^{[l+1]})$$



$$z^{[l+1]} = W^{[l+1]} a^{[l]} + b^{[l+1]}$$

$$z_{\text{norm}}^{[l+1]} = \frac{z^{[l+1]} - \mu_z^{[l+1]}}{\sigma_z^{[l+1]}}$$

$$\tilde{z}^{[l+1]}$$

$$\tilde{z}^{[l+1]} = \gamma^{[l+1]} \odot z_{\text{norm}}^{[l+1]} + \beta^{[l+1]}$$

↑  
Vector. element-wise

$$a^{[l+1]} = \sigma(\tilde{z}^{[l+1]})$$

# Training with Batch Norm.

- Original:  $\{\underline{w}^{[1]}, \underline{b}^{[1]}, \dots, \underline{w}^{[L]}, \underline{b}^{[L]}\}$

- With Batch Norm.:  $\{\underline{w}^{[1]}, \underline{b}^{[1]}, \beta^{[1]}, \gamma^{[1]}, \dots, \underline{w}^{[L]}, \underline{b}^{[L]}, \beta^{[L]}, \gamma^{[L]}\}$

$$w, b, \beta, \gamma \leftarrow w, b, \beta, \gamma - \eta \nabla \mathcal{L}(w, b, \beta, \gamma)$$

learning

with mini-Batch

$B = 8, 16, 32$

$$X^{\{1\}}, X^{\{2\}}, X^{\{3\}}, \dots, X^{\{m\}} = \{X^{\{m\}}(1), \dots, X^{\{m\}}(B)\}$$

$$\mu^{[L]} = \frac{1}{B} \sum_{b=1}^B \sum_{i=1}^n x^{[L]\{m\}}(b)_i$$

$$\sigma^{[L]} = \frac{1}{B-1} \sum_{b=1}^B \left( \sum_{i=1}^n (x^{[L]\{m\}}(b)_i - \mu^{[L]})^2 \right)$$

# Batch Norm at test time

$$\underline{\underline{\mu, \sigma^{[l]}}}$$

$$\mu^{[l]} \approx \hat{\mu}^{[l]} \leftarrow \text{last mini-batch.}$$

$$\approx \underline{\underline{\hat{\mu}^{[l]}}} \leftarrow \text{using all training data.}$$

$$\textcircled{\hat{\mu}^{[l]}} = \frac{\eta^M \hat{\mu}^{[l]\{1\}} + \eta^{M-1} \hat{\mu}^{[l]\{2\}} + \dots + \eta \hat{\mu}^{[l]\{M\}}}{\eta^M + \dots + \eta} \quad \eta < 1$$

$$\Rightarrow \boxed{\hat{\mu}^{[l]}} \leftarrow \textcircled{(1-\eta)} \boxed{\hat{\mu}^{[l]\{M\}}} + \textcircled{\eta} \boxed{\hat{\mu}^{[l]}}$$

# Batch Norm as regularization

- Each mini-batch is scaled by the mean/variance computed on just that mini-batch.
- This adds some noise to the values  $z^{[l]}$  within that minibatch. So similar to dropout, it adds some noise to each hidden layer's activations.
- This has a slight regularization effect.